

# DGM assignment 1 solution

---

## Question 1

### 1.1 Log-Likelihood Function

Maximum Likelihood Estimation (MLE) of an i.i.d. sample:

Given an i.i.d. sample  $(X_t; t = 1, \dots, T)$  where  $X_t \sim f_\theta(x)$ , the likelihood function is:

$$L(\theta) = f_\theta(X_1, X_2, \dots, X_T) = \prod_{t=1}^T f_\theta(X_t),$$

due to independence. Therefore, the log-likelihood function is:

$$\log L(\theta) = \sum_{t=1}^T \log f_\theta(X_t),$$

and the maximum likelihood estimate is:

$$\hat{\theta}_{ML} = \arg \max_{\theta} \log L(\theta).$$

For the AR(1) model:

$$Y_t = \phi Y_{t-1} + \varepsilon_t, \quad t = 1, \dots, T,$$

where  $\varepsilon_t \sim \text{i.i.d. } N(0, \sigma^2)$  and  $y_0$  is fixed/known. The parameter vector is  $\Theta = (\phi, \sigma^2)$ .

The joint probability can be expressed using the chain rule:

$$f_\theta(y_0, y_1, \dots, y_T) = f_\theta(y_T | y_{T-1}, \dots, y_0) f_\theta(y_{T-1}, \dots, y_0).$$

By recursively applying this, we get:

$$f_\theta(y_0, y_1, \dots, y_T) = \prod_{t=1}^T f_\theta(y_t | y_{t-1}, \dots, y_0) f_\theta(y_0).$$

Since  $Y_t$  given past values depends only on  $Y_{t-1}$ :

$$Y_t | Y_{t-1}, \dots, y_0 \equiv Y_t | Y_{t-1},$$

so:

$$f_\theta(Y_t | Y_{t-1}, \dots, y_0) = f_\theta(Y_t | Y_{t-1}).$$

The conditional distribution is:

$$Y_t | Y_{t-1} \sim N(E(Y_t | Y_{t-1}), V(Y_t | Y_{t-1})),$$

where:

$$E(Y_t | Y_{t-1}) = \phi Y_{t-1}, \quad V(Y_t | Y_{t-1}) = \sigma^2.$$

Thus, the conditional density is:

$$f_\theta(Y_t | Y_{t-1}) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(Y_t - \phi Y_{t-1})^2}{2\sigma^2}\right).$$

The conditional likelihood function is:

$$L(\theta) = f_\theta(Y_1, Y_2, \dots, Y_T | y_0) = \prod_{t=1}^T f_\theta(Y_t | Y_{t-1}).$$

Therefore, the log-likelihood function is:

$$\log L(\theta) = \sum_{t=1}^T \log f_\theta(Y_t | Y_{t-1}).$$

## 1.2 Maximum Likelihood Estimation

To find the maximum likelihood estimates, we solve:

$$\hat{\theta} = \arg \max_{\theta} \log L(\theta).$$

The estimate for  $\phi$  is:

$$\hat{\phi} = \frac{\frac{1}{T} \sum_{t=1}^T Y_{t-1} Y_t}{\frac{1}{T} \sum_{t=1}^T Y_{t-1}^2},$$

which is the Ordinary Least Squares (OLS) estimator.

The estimate for  $\sigma^2$  is:

$$\hat{\sigma}^2 = \frac{1}{T} \sum_{t=1}^T \hat{\epsilon}_t^2, \quad \hat{\epsilon}_t = Y_t - \hat{\phi} Y_{t-1}.$$

---

## Question 2

### 2.1 Maximum Likelihood Estimation

Deriving the Log-Likelihood Function:

The joint probability distribution of the sequence  $x = (x_1, x_2, \dots, x_T)$  is given by:

$$p(x) = p(x_1) \prod_{t=2}^T p(x_t | x_1, x_2, \dots, x_{t-1})$$

Given that each conditional distribution  $p(x_t | x_1, x_2, \dots, x_{t-1})$  is modeled as a Gaussian with mean  $\mu_t = f_{\theta}(x_1, x_2, \dots, x_{t-1})$  and fixed variance  $\sigma^2$ :

$$p(x_t | x_1, \dots, x_{t-1}) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_t - \mu_t)^2}{2\sigma^2}\right)$$

The log-likelihood  $L(\theta)$  of the joint distribution  $p(x)$  is

$$L(\theta) = \log p(x) = \sum_{t=1}^T \log p(x_t | x_1, \dots, x_{t-1})$$

Substituting the expression for the Gaussian conditional distributions:

$$L(\theta) = \sum_{t=1}^T \left( -\frac{1}{2} \log(2\pi\sigma^2) - \frac{(x_t - f_{\theta}(x_1, \dots, x_{t-1}))^2}{2\sigma^2} \right)$$

Simplifying constants:

$$L(\theta) = -\frac{T}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{t=1}^T (x_t - f_{\theta}(x_1, \dots, x_{t-1}))^2$$

Computing Gradients Using Backpropagation:

To train the neural network  $f_{\theta}$ , we aim to maximize the log-likelihood  $L(\theta)$  or equivalently minimize the negative log-likelihood (NLL):

$$NLL(\theta) = -L(\theta) = \frac{1}{2\sigma^2} \sum_{t=1}^T (x_t - f_{\theta}(x_1, \dots, x_{t-1}))^2 + \text{const}$$

This is equivalent to the mean squared error (MSE) loss scaled by  $1/(2\sigma^2)$ . The gradient of the NLL with respect to  $\theta$  is:

$$\nabla_{\theta} NLL(\theta) = -\frac{1}{\sigma^2} \sum_{t=1}^T (x_t - f_{\theta}(x_1, \dots, x_{t-1})) \nabla_{\theta} f_{\theta}(x_1, \dots, x_{t-1})$$

Backpropagation Process:

### 1. Forward Pass:

- For each time step  $t$ :
  - Input  $x_1, \dots, x_{t-1}$  into the neural network  $f_\theta$  to compute  $\mu_t = f_\theta(x_1, \dots, x_{t-1})$ .
  - Compute the prediction error  $e_t = x_t - \mu_t$ .

### 2. Loss Computation:

- Accumulate the loss over all time steps:

$$\text{Loss} = \frac{1}{2\sigma^2} \sum_{t=1}^T e_t^2$$

### 3. Backward Pass:

- Backpropagate the gradient of the loss with respect to the network parameters  $\theta$  through  $f_\theta$ :

$$\nabla_\theta \text{Loss} = -\frac{1}{\sigma^2} \sum_{t=1}^T e_t \nabla_\theta f_\theta(x_1, \dots, x_{t-1})$$

- Update  $\theta$  using gradient descent or an appropriate optimization algorithm.

## 2.2 Predictive Sampling

Sampling Process:

To generate a new sequence  $x' = (x'_1, x'_2, \dots, x'_T)$ :

#### 1. Initialize $x'_1$ :

- If  $p(x_1)$  is specified, sample  $x'_1 \sim p(x_1)$ .
- If not specified, assume  $x'_1 \sim N(\mu_0, \sigma^2)$ , where  $\mu_0 = f_\theta()$ .

#### 2. Iterative Sampling for $t = 2$ to $T$ :

- Compute Mean:

$$\mu_t = f_\theta(x'_1, x'_2, \dots, x'_{t-1})$$

- Sample  $x'_t$ :

$$x'_t \sim N(\mu_t, \sigma^2)$$

Influence of Autoregressive Structure:

- Each  $x'_t$  depends on all previous samples  $x'_1, \dots, x'_{t-1}$ , making the sampling process sequential.
- The autoregressive model captures temporal dependencies by conditioning  $x'_t$  on the entire history up to  $t - 1$ .
- The sequential nature ensures that long-term dependencies can, in theory, influence future samples.

## 2.3 KL Divergence between AR Models

### 1. KL Divergence Definition:

Given two probability distributions,  $p_\theta(x)$  and  $q_\phi(x)$ , parameterized by two different models ( $f_\theta$  and  $g_\phi$ ), the KL divergence is defined as:

$$D_{KL}(p_\theta(x) \parallel q_\phi(x)) = E_{x \sim p_\theta} \left[ \log \frac{p_\theta(x)}{q_\phi(x)} \right]$$

This measures the "distance" between the two distributions  $p_\theta(x)$  and  $q_\phi(x)$ . It quantifies how much information is lost when using  $q_\phi$  to approximate  $p_\theta$ .

### 2. Autoregressive Models Factorization:

For autoregressive models, the joint probability of the sequence  $x = (x_1, \dots, x_T)$  is factorized as:

- Model 1:

$$p_\theta(x) = p_\theta(x_1) \prod_{t=2}^T p_\theta(x_t | x_1, \dots, x_{t-1})$$

where  $p_\theta(x_t | x_1, \dots, x_{t-1})$  is modeled by the neural network  $f_\theta$ .

- Model 2:

$$q_\phi(x) = q_\phi(x_1) \prod_{t=2}^T q_\phi(x_t | x_1, \dots, x_{t-1})$$

where  $q_\phi(x_t | x_1, \dots, x_{t-1})$  is modeled by the neural network  $g_\phi$ .

### 3. KL Divergence for Autoregressive Models:

Using the factorized forms of  $p_\theta(x)$  and  $q_\phi(x)$ , the KL divergence between them can be written as:

$$D_{KL}(p_\theta(x) \parallel q_\phi(x)) = E_{x \sim p_\theta} \left[ \log \frac{p_\theta(x_1)}{q_\phi(x_1)} + \sum_{t=2}^T \log \frac{p_\theta(x_t | x_1, \dots, x_{t-1})}{q_\phi(x_t | x_1, \dots, x_{t-1})} \right]$$

### 4. KL Divergence for Gaussian Distributions:

Since both models  $p_\theta(x_t | x_1, \dots, x_{t-1})$  and  $q_\phi(x_t | x_1, \dots, x_{t-1})$  are Gaussian:

$$p_\theta(x_t | x_1, \dots, x_{t-1}) = \mathcal{N}(f_\theta(x_1, \dots, x_{t-1}), \sigma_\theta^2)$$

$$q_\phi(x_t | x_1, \dots, x_{t-1}) = \mathcal{N}(g_\phi(x_1, \dots, x_{t-1}), \sigma_\phi^2)$$

The KL divergence between two univariate Gaussians  $\mathcal{N}(\mu_1, \sigma_1^2)$  and  $\mathcal{N}(\mu_2, \sigma_2^2)$  is given by:

$$D_{KL}(\mathcal{N}(\mu_1, \sigma_1^2) \parallel \mathcal{N}(\mu_2, \sigma_2^2)) = \log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2}$$

Thus, for each time step  $t$ , the KL divergence between the conditional distributions is:

$$KL(p_\theta(x_t | x_1, \dots, x_{t-1}) \parallel q_\phi(x_t | x_1, \dots, x_{t-1})) = \log \frac{\sigma_\phi}{\sigma_\theta} + \frac{\sigma_\theta^2 + (f_\theta(x_1, \dots, x_{t-1}) - g_\phi(x_1, \dots, x_{t-1}))^2}{2\sigma_\phi^2} - \frac{1}{2}$$

### 5. Full KL Divergence Expression:

Now, summing over all time steps, the total KL divergence between the two autoregressive models is:

$$D_{KL}(p_\theta(x) \parallel q_\phi(x)) = \sum_{t=1}^T E_{x \sim p_\theta} \left[ \log \frac{\sigma_\phi}{\sigma_\theta} + \frac{\sigma_\theta^2 + (f_\theta(x_1, \dots, x_{t-1}) - g_\phi(x_1, \dots, x_{t-1}))^2}{2\sigma_\phi^2} - \frac{1}{2} \right]$$

### 6. Computational Challenges:

Directly evaluating the KL divergence for high-dimensional sequences presents the following challenges:

- **Expectation over the distribution  $p_\theta$ :** We need to compute the expectation  $E_{x \sim p_\theta}$ , which requires generating samples from the autoregressive model  $p_\theta(x)$  and averaging the KL divergence over those samples.

- **Sequential Dependence:** In autoregressive models, each sample  $x_t$  depends on all previous values, which means generating each sample requires sequential computation. This can become computationally expensive for long sequences.

### 7. Approximation Method:

A practical approach to approximate the KL divergence is to use Monte Carlo sampling:

- Generate  $K$  samples  $\{x^{(i)}\}_{i=1}^K$  from the distribution  $p_\theta(x)$ .

- Estimate the KL divergence as the empirical average over these samples:

$$D_{KL}(p_\theta(x) \parallel q_\phi(x)) \approx \frac{1}{K} \sum_{i=1}^K \left[ \log \frac{p_\theta(x^{(i)})}{q_\phi(x^{(i)})} \right]$$

- For each sampled sequence  $x^{(i)}$ , compute the KL divergence between the two models' outputs at each time step and average over the samples.

## 2.4 Stationarity and Long-Term Dependencies

Analysis of Long-Term Dependencies:

- **Model Capacity:** The autoregressive model conditions  $x_t$  on all previous observations  $x_{<t}$ , which theoretically allows capturing long-term dependencies.

- **Practical Limitations:** - **Computational Complexity:** Feeding all past observations into  $f_\theta$  becomes impractical as  $t$  increases. - **Vanishing/Exploding Gradients:** Standard feedforward networks struggle with learning long-term dependencies due to gradient issues during backpropagation. - **Fixed Input Size:** Neural networks with fixed-size inputs cannot handle varying-length histories efficiently.

Mathematical Explanation:

- **Inefficient Memory:** The model does not prioritize relevant past information, treating all past observations equally, which dilutes the influence of distant dependencies.

- **Curse of Dimensionality:** As  $t$  grows, the input space becomes vast, making learning and generalization difficult.

Suggested Modification: Using Recurrent Neural Networks (RNNs) or Transformers

**Recurrent Neural Networks:**

- **Architecture:** - Introduce a hidden state  $h_t$  that summarizes past information:

$$h_t = \text{RNNCell}(h_{t-1}, x_{t-1})$$

- The conditional mean becomes:

$$\mu_t = f_\theta(h_t)$$

**Advantages:**

- **Sequential Processing:** Efficiently processes sequences of arbitrary length.
- **Parameter Sharing:** Reuses weights across time steps, reducing the number of parameters.
- **Capturing Long-Term Dependencies:** Designed to maintain information over time, especially with variants like LSTMs or GRUs.

**Transformers:**

- **Architecture:** - Uses self-attention mechanisms to weigh the importance of past observations:

$$\mu_t = f_\theta(x_{<t}) = \text{Transformer}(x_{<t})$$

- **Advantages:** - **Parallelization:** Processes all positions simultaneously, improving computation speed. - **Attention Mechanism:** Explicitly models dependencies between all pairs of positions, effectively capturing long-term dependencies. - **Positional Encoding:** Incorporates the order of the sequence, essential for temporal data.

Probabilistic Interpretation:

- **Hidden States as Sufficient Statistics:** - The hidden state  $h_t$  or attention weights act as sufficient statistics summarizing past information needed to predict  $x_t$ .

- **Modified Conditional Distribution:**

$$p(x_t|h_t) = \mathcal{N}(f_\theta(h_t), \sigma^2)$$

- **Efficient Dependency Modeling:** - By summarizing past information, the model focuses on relevant dependencies, enhancing its ability to capture long-term patterns.

## Question 3

autoregressive model

$$p(x) = p(x_1)p(x_2|x_{<2}) \dots p(x_i|x_{<i}) \dots p(x_n|x_{<n}),$$

and Real NADE models the conditional distribution as

$$p(x_1) = \mathcal{N}(x_1|\mu_1, \exp(s_1)),$$

$$\vdots$$

$$p(x_i|x_{<i}; W, c, v_i, b_i, u_i, d_i) = \mathcal{N}\left(x_i \middle| v_i^\top h_i + b_i, \exp(u_i^\top h_i + d_i)\right),$$

where  $h_i, c, v_i, u_i \in R^d$ .

Now, we would like to make  $p(x_i|x_{<i})$  follow a mixture of Gaussian

$$p(x_i|x_{<i}) = \sum_{c=1}^C \pi_i^c \mathcal{N}(\mu_i^c, (\sigma_i^c)^2),$$

where  $\sum_{c=1}^C \pi_i^c = 1$ .

Now the question is: How do you propose to parameterize  $\pi_i^c, \mu_i^c, \sigma_i^c, \forall c \in \{1, \dots, C\}$  as a function of  $h_i$ ? Describe the parameters required and the total number of parameters required for a single  $p(x_i|x_{<i})$ .

## Answer

We can let

$$\pi_i^c = \frac{\exp((g_i^c)^\top h_i)}{\sum_{c'=1}^C \exp((g_i^{c'})^\top h_i)}, \quad \forall c \in \{1, \dots, C\}$$

where  $g_i^c \in R^d$ .

In addition, we have

$$\mu_i^c = (v_i^c)^\top h_i, \quad \sigma_i^c = \exp((u_i^c)^\top h_i), \quad \forall c \in \{1, \dots, C\}$$

where  $h_i \in R^d$ ,  $v_i^c \in R^d$ , and  $u_i^c \in R^d$ .

Therefore, the parameters required for each  $p(x_i|x_{<i})$  are:

- For the mixture weights  $\pi_i^c$ : parameters  $g_i^c \in R^d$  for each component  $c$ , totaling  $C \times d$  parameters.
- For the means  $\mu_i^c$ : parameters  $v_i^c \in R^d$  for each component  $c$ , totaling  $C \times d$  parameters.
- For the standard deviations  $\sigma_i^c$ : parameters  $u_i^c \in R^d$  for each component  $c$ , totaling  $C \times d$  parameters.

Thus, the total number of parameters required for a single  $p(x_i|x_{<i})$  is:

$$\text{Total Parameters} = 3 \times C \times d$$

## Question 4

### 4.1 A Quick Warm Up

What is the value of this expectation  $E_{x \sim N(0,2)}[x^2 + x + 1]$ ? How do you propose to estimate this quantity with Monte Carlo estimation?

## Answer

We know that for a Gaussian distribution  $N(0, 2)$ :

$$E[x] = 0, \quad E[x^2] = \text{Var}(x) = 2.$$

So we can compute:

$$E_{x \sim N(0,2)}[x^2 + x + 1] = E[x^2] + E[x] + 1 = 2 + 0 + 1 = 3.$$

To estimate this value using Monte Carlo estimation, we can:

1. Randomly sample  $K$  values  $x_i \sim N(0, 2)$  for  $i = 1, 2, \dots, K$ .
2. Compute  $y_i = x_i^2 + x_i + 1$  for each sample.
3. Estimate the expectation as:

$$\hat{y} = \frac{1}{K} \sum_{i=1}^K y_i.$$

## 4.2 Variance of K-sample Estimator

What is the variance of the Monte Carlo estimator using  $K$  samples?

### Answer

Since each  $y_i$  is an independent and identically distributed (i.i.d.) random variable, the variance of the estimator  $\hat{y}$  is:

$$\text{Var}[\hat{y}] = \frac{1}{K} \text{Var}[y_1].$$

First, compute  $\text{Var}[y_1]$ :

$$\text{Var}[y_1] = E[(y_1 - E[y_1])^2] = E[(x^2 + x + 1 - 3)^2] = E[(x^2 + x - 2)^2].$$

Expand the squared term:

$$(x^2 + x - 2)^2 = x^4 + 2x^3 - 3x^2 - 4x + 4.$$

Compute each expected value:

$$\begin{aligned} E[x^4] &= 3 \times (\text{Var}(x))^2 = 3 \times 2^2 = 12, \\ E[x^3] &= 0, \quad \text{since } x \text{ is symmetric around zero,} \\ E[x^2] &= 2, \\ E[x] &= 0. \end{aligned}$$

Now, compute the variance:

$$\text{Var}[y_1] = 12 + 0 - 3 \times 2 - 0 + 4 = 12 - 6 + 4 = 10.$$

Therefore, the variance of the estimator is:

$$\text{Var}[\hat{y}] = \frac{1}{K} \times 10 = \frac{10}{K}.$$

## 4.3 Objective Minimization

Now assume we are interested in minimizing the objective:

$$F(\theta) = \sum_{n=1}^N w_n f(\theta; n) + \lambda R(\theta), \quad \text{where } w_n > 0, \forall n.$$

What is the asymptotic complexity of evaluating  $F(\theta)$  and  $\nabla F(\theta)$  given  $\theta$  (assume the function call and gradient evaluation of  $f(\theta; n)$  is constant)? How do you propose to use Monte Carlo estimation to acquire an unbiased estimation of the objective and its gradient?

### Answer

#### Asymptotic Complexity:

- Evaluating  $F(\theta)$  requires summing over  $N$  terms, so the complexity is  $O(N)$ . - Computing  $\nabla F(\theta)$  also requires summing over  $N$  gradients  $\nabla f(\theta; n)$ , so the complexity is  $O(N)$ .

#### Monte Carlo Estimation:

To reduce computational complexity, we can use Monte Carlo estimation to approximate  $F(\theta)$  and  $\nabla F(\theta)$ .

#### 1. Normalize Weights:

Define  $W = \sum_{n=1}^N w_n$  and probability distribution  $p(n) = \frac{w_n}{W}$ .

#### 2. Sample Indices:

Sample  $K$  indices  $\{n_i\}_{i=1}^K$  from the categorical distribution  $p(n)$ .

### 3. Estimate Objective:

The unbiased estimator of  $F(\theta)$  is:

$$\hat{F}(\theta) = W \left( \frac{1}{K} \sum_{i=1}^K f(\theta; n_i) \right) + \lambda R(\theta).$$

### 4. Estimate Gradient:

The unbiased estimator of  $\nabla F(\theta)$  is:

$$\nabla \hat{F}(\theta) = W \left( \frac{1}{K} \sum_{i=1}^K \nabla f(\theta; n_i) \right) + \lambda \nabla R(\theta).$$

By using a subset of  $K$  samples (where  $K \ll N$ ), we reduce the computational complexity to  $O(K)$ , which is more efficient for large  $N$ .

---